

Advanced Parallel Programming: Understanding Locking and Shared Variables



Filip Ekberg

Principal Consultant & CEO

@fekberg fekberg.com



Working with Shared Variables



ConcurrentBag<T>

```
var bag = new ConcurrentBag<string>();
```

```
Parallel.Invoke(  
    () => { bag.Add( "Hello" ); },  
    () => { bag.Add( "World" ); },  
    () => { bag.Add( "from" ); },  
    () => { bag.Add( "Pluralsight" ); }  
);
```



Normal For vs Parallel.For

```
for(int i = 0; i < 100; i++)  
{  
    // Execute sequentially  
}
```

```
Parallel.For(0, 100, (i) => {  
    // Execute in parallel  
});
```



Be careful when **sharing**
variables and **resources** in a
parallel process as you will
run into **race conditions**



**Always prefer atomic
operations** over lock when
possible as its **less overhead**
and performs **faster**



Introducing a Lock

```
Parallel.For(0, 100, (i) => {  
    lock(...)  
    {  
    }  
});
```



Introducing a Lock

```
Parallel.For(0, 100, (i) => {  
    lock(...)  
    {  
        // Only one thread at a time please..  
    }  
});
```



Introducing a Lock

```
Parallel.For(0, 100, (i) => {  
    lock(syncRoot)  
    {  
  
    }  
});
```



**Lock object is often
referred to as sync root**



Introducing a Lock

```
Parallel.For(0, 100, (i) => {  
    lock(syncRoot)  
    {  
    }  
});
```



If another thread has locked the “door” the thread will be blocked until the “door” is opened



Only **lock** for as **short** of a
time as possible



Use a **lock** to **access**
resources in a **thread-safe**
manner



Considerations



Be careful when adding a lock



Nested locks can lead to deadlocks



Only lock in short periods of time



Do as little work as possible in the lock statement



Next: Performing Atomic Operations



Performing Atomic Operations



Interlocked from System.Threading

“Provides atomic operations for variables that are shared by multiple threads.”



Using Interlocked

```
Interlocked.Increment(ref int thisValue)
```

```
Interlocked.Decrement(ref int thisValue)
```

```
Interlocked.Add(ref int toBeUpdated, int withValue)
```

```
Interlocked.Add(ref long toBeUpdated, long withValue)
```



Incrementing a Variable

```
int result = 0;  
Parallel.For(0, 100, (i) => {  
  
});
```



Incrementing a Variable

```
static object syncRoot = new object();

int result = 0;
Parallel.For(0, 100, (i) => {
    lock(syncRoot)
    {
        result += 1;
    }
});
```



Interlocked.Increment

```
int result = 0;  
Parallel.For(0, 100, (i) => {  
    Interlocked.Increment(ref result)  
});
```



Interlocked.Increment

```
int result = 0;  
Parallel.For(0, 100, (i) => {  
    Interlocked.Increment(ref result)  
});
```



Only one thread performs the read & write at a given time. This is faster than using a lock()



Interlocked vs Lock

Interlocked

```
int result = 0;
Parallel.For(0, 100, (i) => {
    Interlocked.Increment(ref result)
});
```

Lock

```
static object syncRoot =
    new object();

int result = 0;
Parallel.For(0, 100, (i) => {
    lock(syncRoot)
    {
        result += 1;
    }
});
```

Interlocked.Decrement

```
int result = 0;  
Parallel.For(0, 100, (i) => {  
    Interlocked.Decrement(ref result)  
});
```



Interlocked uses **less instructions** and is **faster than a lock**



Bitwise Operations with Interlocked

```
int Interlocked.Or(ref int thisValue, int withThis)  
int Interlocked.And(ref int thisValue, int withThis)
```



Returns the original value of “thisValue”



Set a New Value

```
T Interlocked.Exchange<T>(ref T source, T newValue)
```



Set a New Value

```
T Interlocked.Exchange<T>(ref T source, T newValue)
```



Returns the “source” value



Interlocked.Exchange

```
int Interlocked.Exchange(ref int source, int newValue)
```



Exercise: Use Interlocked.Exchange as a Gate Keep

```
var gate = 0;

// Try locking the gate
var state = Interlocked.Exchange(ref gate, 1);

if(state == 0)
{
    // Gate was previously locked
}
else
{
    // Gate was already locked

    // Back-off for a while
    // Retry
}
```



Next: Deadlocks with Nested Locks



Deadlocks with Nested Locks



Locked? Thread Will Be Blocked Until Opened

```
Task.Run(() => {  
    lock(syncRoot)  
    {  
        Thread.Sleep(5000);  
    }  
});
```

```
Task.Run(() => {  
    lock(syncRoot)  
    {  
  
    }  
});
```



Locked? Thread Will Be Blocked Until Opened


```
Task.Run(() => {  
    lock(syncRoot)  
    {  
        Thread.Sleep(5000);  
    }  
});
```

```
Task.Run(() => {  
    lock(syncRoot)  
    {  
    }  
});
```

**This thread will be locked
until “syncRoot” is released (5 seconds)**




Deadlock: Fighting over the Same Resource



Thread 1


Waiting for **Thread 2** to complete to
access **Resource 2**




Thread 2

Waiting for **Thread 1** to complete to
access **Resource 1**

Avoiding a Deadlock




Don't share lock objects for multiple shared resources



Don't use a **string** as a lock




Use one lock object for each shared resource



Don't use a type instance from **typeof()** as a lock



Give you lock object a meaningful name



Don't use “**this**” as a lock



Nested locks may result in a
deadlock



Can You Guarantee That This Doesn't Deadlock?

```
lock(lock1)
{
    Parallel.For(0, 100, (i) => {
        Deadlock();
    });
}
```



Can You Guarantee That This Doesn't Deadlock?

```
lock(lock1)
{
    Parallel.For(0, 100, (i) => {
        Deadlock();
    });
}
```

```
void Deadlock()
{
    lock(lock1)
    {
        Thread.Sleep(1);
    }
}
```



**Avoid nested locks and
shared locks!**



Next: Cancel Parallel Operations



Cancel Parallel Operations



Cancellation

```
CancellationTokenSource cts = new CancellationTokenSource();

var options = new ParallelOptions
{
    CancellationToken = cts.Token
};

Parallel.For(0, 100, options, _ => {});

Parallel.ForEach(source, options, _ => {});

Parallel.Invoke(options, () => {}, () => {});
```



When a **cancellation** is
detected: No further
iterations/operations will
start!



Parallel Methods

**Built on the Task
Parallel Library**

**Monitors for
Cancellation**

**A task doesn't start
when a
cancellation token
is cancelled**



It **won't stop** executing
already **started operations**



Monitor the Cancellation Token

```
CancellationTokenSource cts = new CancellationTokenSource();

var options = new ParallelOptions
{
    CancellationToken = cts.Token
};

Parallel.For(0, 100, options, _ => {
    if(cts.Token.IsCancellationRequested)
    {
        // Roll back?
    }
});
```



When to Monitor the Cancellation Token

```
Parallel.For(0, 100, options, _ => {  
  
    FirstExpensiveOperation();  
  
    if(cts.Token.IsCancellationRequested)  
    {  
  
    }  
    else  
    {  
        SecondExpensiveOperation();  
    }  
});
```



Each system executing the
code **might complete** it
faster, or slower!



ThreadLocal and AsyncLocal Variables



ThreadLocal<T>

Provides storage that is local to a thread

Note: Task in the Task Parallel Library **reuse threads!**



Thread Local/Static data
may be **shared between**
multiple tasks that uses the
same thread



ThreadLocal<T>

```
ThreadLocal<decimal?> data = new ThreadLocal<decimal?>();
```



ThreadLocal<T>

```
ThreadLocal<decimal?> data = new ThreadLocal<decimal?>();
```

```
data.Value = 100m;
```



ThreadLocal<T>

```
ThreadLocal<decimal?> data = new ThreadLocal<decimal?>();
```

```
data.Value = 100m;
```

```
var values = data.Values; // All values for all threads
```



AsyncLocal<T>

“Represents ambient **data** that is **local** to a **given asynchronous control flow**, such as an **asynchronous method**.”



When you **write** to
AsyncLocal a **local copy** will
be **created**.

The **outer contexts** value will
not be **overwritten**!



AsyncLocal<T>

```
AsyncLocal<decimal?> data = new AsyncLocal<decimal?>();
```

```
data.Value = 100m;
```

```
Task.Run(() => {  
    // Local copy created  
    data.Value = 200m;  
});
```

```
// data.Value is stil 100 after the Task.Run!
```



Summary



When sharing variables and resources be cautious with locking

Don't run expensive operations in a lock

Atomic operations should be preferred when possible as they use less resources, and are faster

Nested locks can cause deadlocks

Don't share lock objects

Cancelling a parallel operation will prevent further operations from starting

ThreadLocal and AsyncLocal are powerful, but very different from each other

